


Lecture 07

Application & Tips:

Learning rate, data preprocessing, overfitting

Gradient descent

Cost (Loss)



Loss (Cost)

$$L = \frac{1}{N} \sum_i D(S(WX_i + b), L_i)$$

Training set

STEP (learning rate)

Gradient descent

```
# Minimize error using cross entropy
```

```
# reduce_sum( , axis=1 ) axis=1 --> row에서 합계를 계산, axis=0 --> column에서 합계를 계산
```

```
Learning_rate = 0.001
```

```
cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
```


```
optimizer = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)
```

=

```
cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
```



```
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.001).minimize(cost)
```

Cost (Loss)




STEP (learning rate)

Large learning rate: overshooting





Small learning rate: takes too long, stops at local minimum



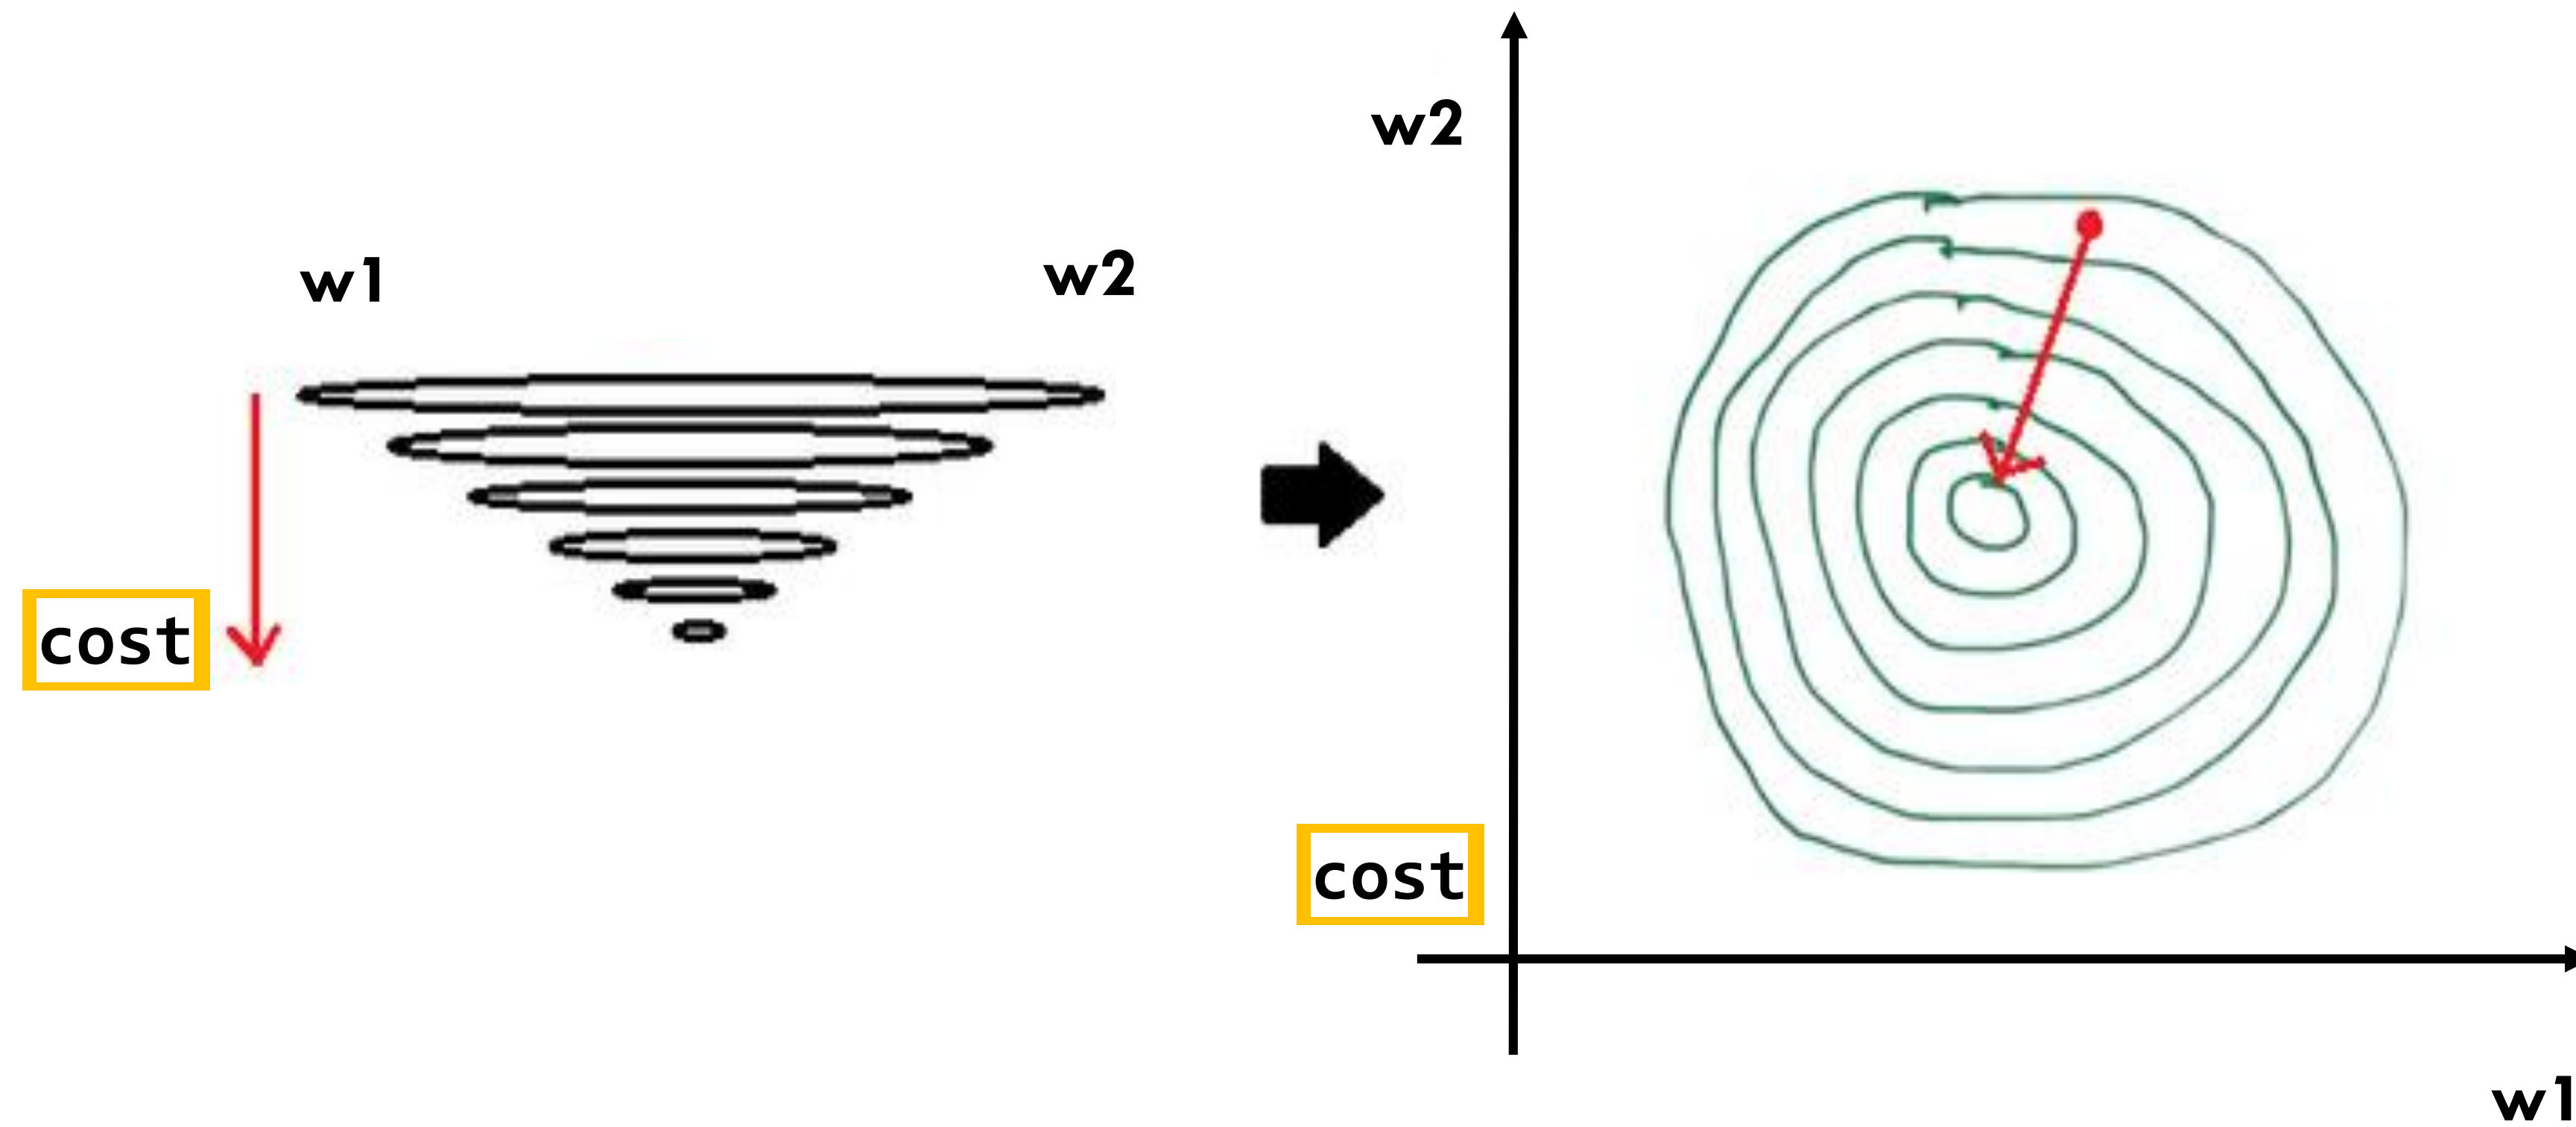
Try several learning rates

- **Observe the cost function**
- **Check it goes down in a reasonable rate**

Data (X) preprocessing for gradient descent




Data (X) preprocessing for gradient descent




Data (X) preprocessing for gradient descent

x1	x2	y
1	9000	A
2	-5000	A
4	-2000	B
6	8000	B
9	9000	C



Data (X) preprocessing for gradient descent



Standardization


$$\mathbf{x}'_j = \frac{\mathbf{x}_j - \mu_j}{\sigma_j}$$

```
X_std[:,0] = (X[:,0] - X[:,0].mean()) / X[:,0].std()
```

Overfitting

- Our model is very good **with training data set (with memorization)**
- **Not good at test dataset or in real use**

Overfitting




Solutions for overfitting

- **More training data!**
- **Reduce the number of features**
- **Regularization**

Regularization

- Let's not have too big numbers in the weight



Loss (Cost)

$$L = \frac{1}{N} \sum_i D(S(WX_i + b), L_i)$$

Training set

Regularization

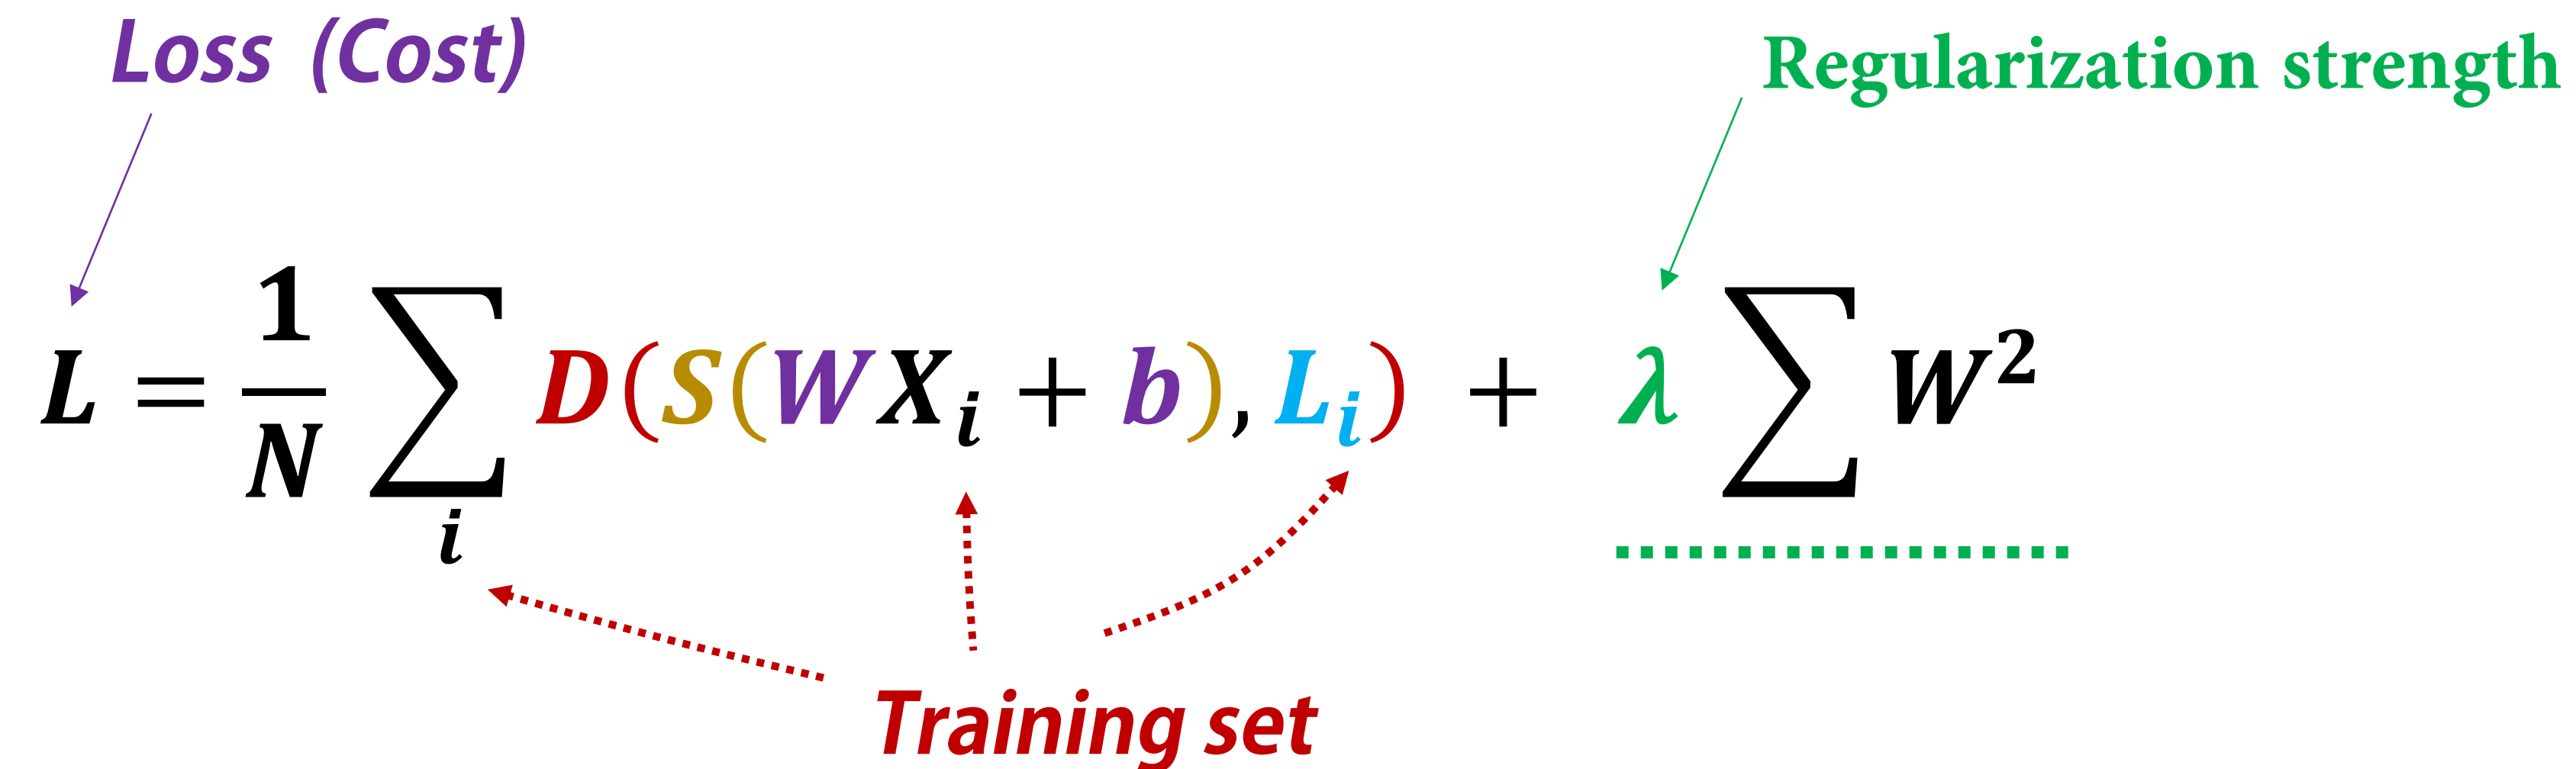
- Let's not have too big numbers in the weight

Loss (Cost)

Regularization strength

$$L = \frac{1}{N} \sum_i D(S(WX_i + b), L_i) + \lambda \sum w^2$$

Training set



Regularization

- Let's not have too big numbers in the weight

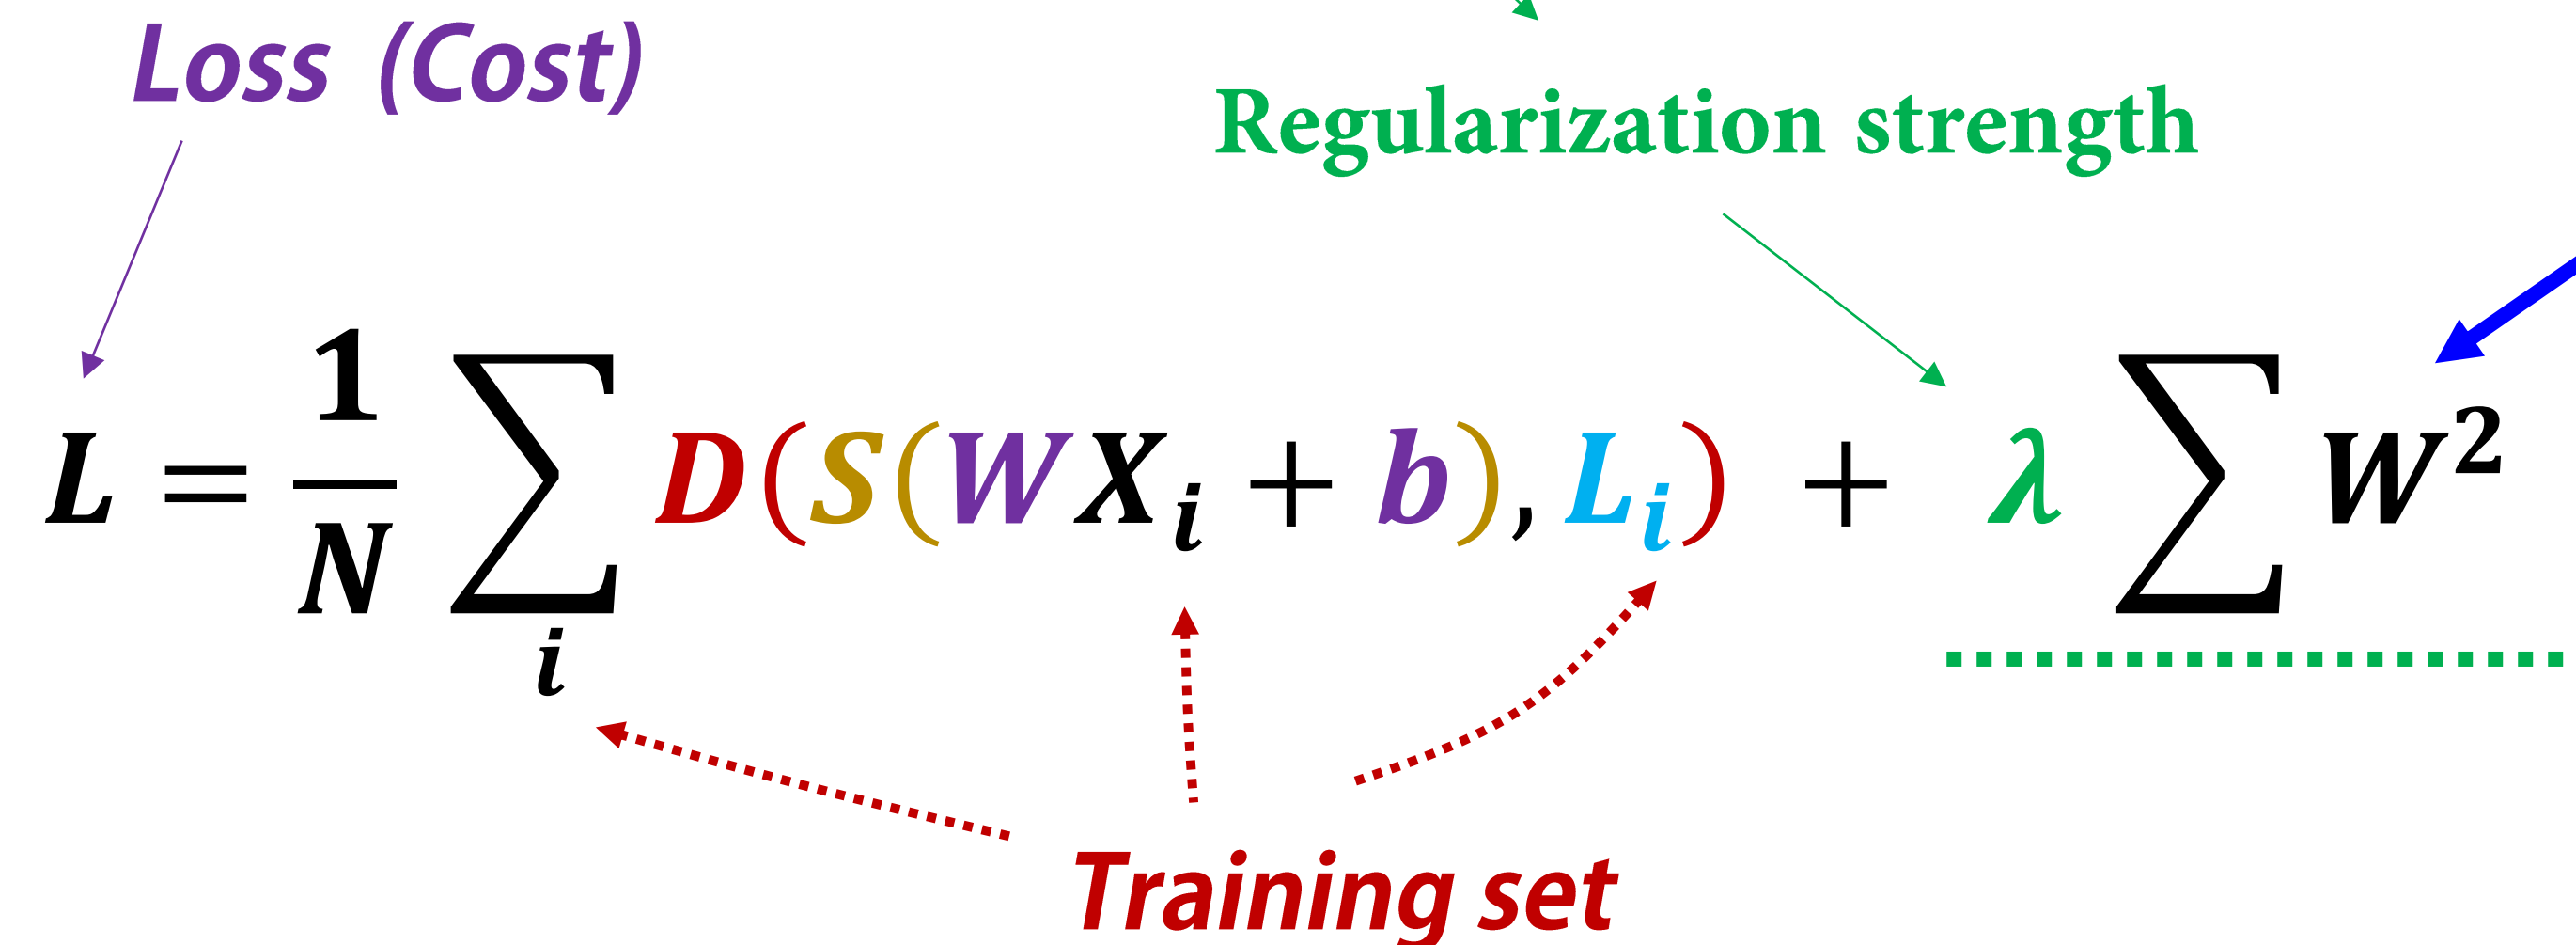
```
l2reg = 0.001 * tf.reduce_sum(tf.square(W))
```

Loss (Cost)

$$L = \frac{1}{N} \sum_i D(S(WX_i + b), L_i) + \lambda \sum w^2$$

Regularization strength

Training set



Summary

- **Learning rate**
- **Data preprocessing**
- **Overfitting**
 - ▷ **More training data**
 - ▷ **Regularization**

Application & Tips:

Learning and test data sets

Performance evaluation: is this good?

data



training

**ML
Model**

Evaluation using training set?

Size	Price
2104	400
1600	330
2400	369
1416	232
3000	540
1985	300
1534	315
1427	199
1380	212
1494	243

training set



**ML
Model**

- **100% correct (accuracy)**
- **Can memorize**

Evaluation using training set?

Size	Price
2104	400
1600	330
2400	369
1416	232
3000	540
1985	300
1534	315

training set



ML
Model

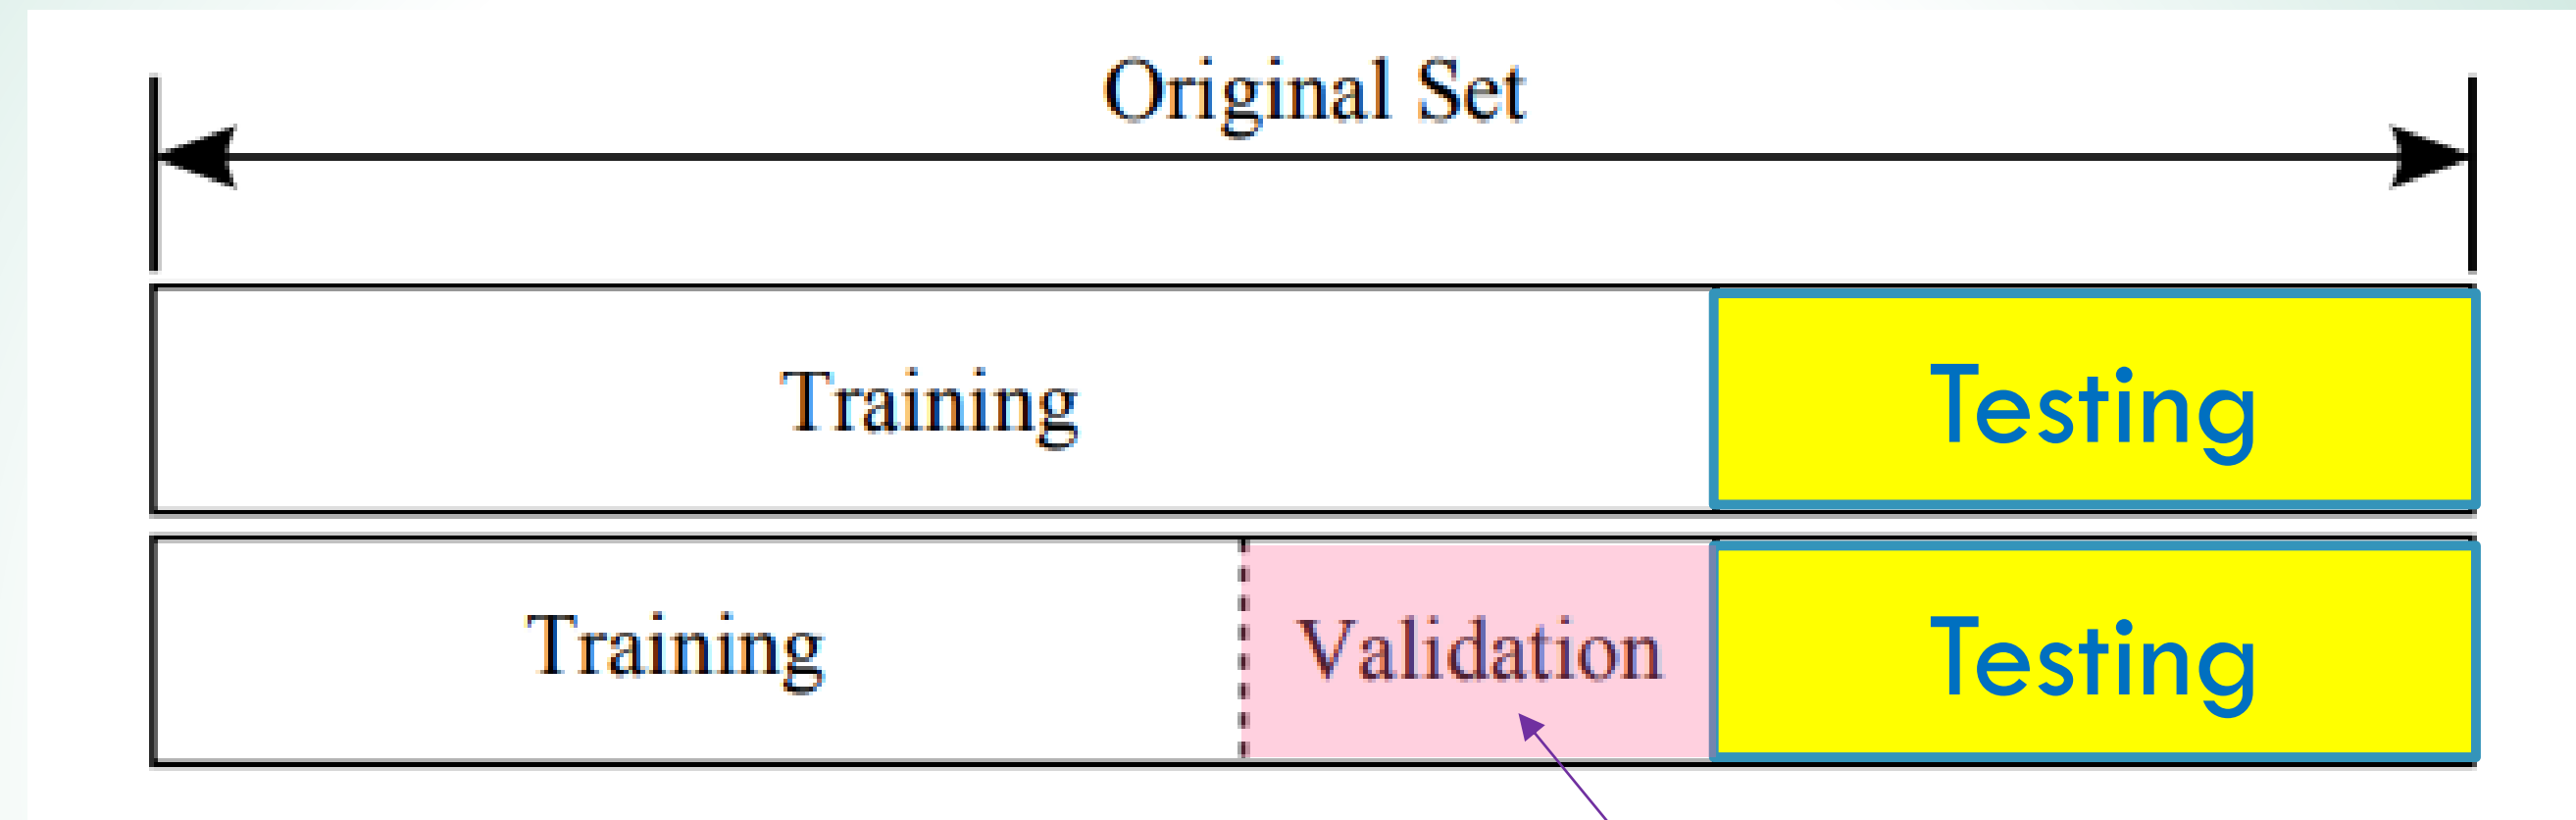
Compare

Y, \hat{Y}

test set

--	--

Training, validation and test sets



Cost (Loss)

Regularization strength

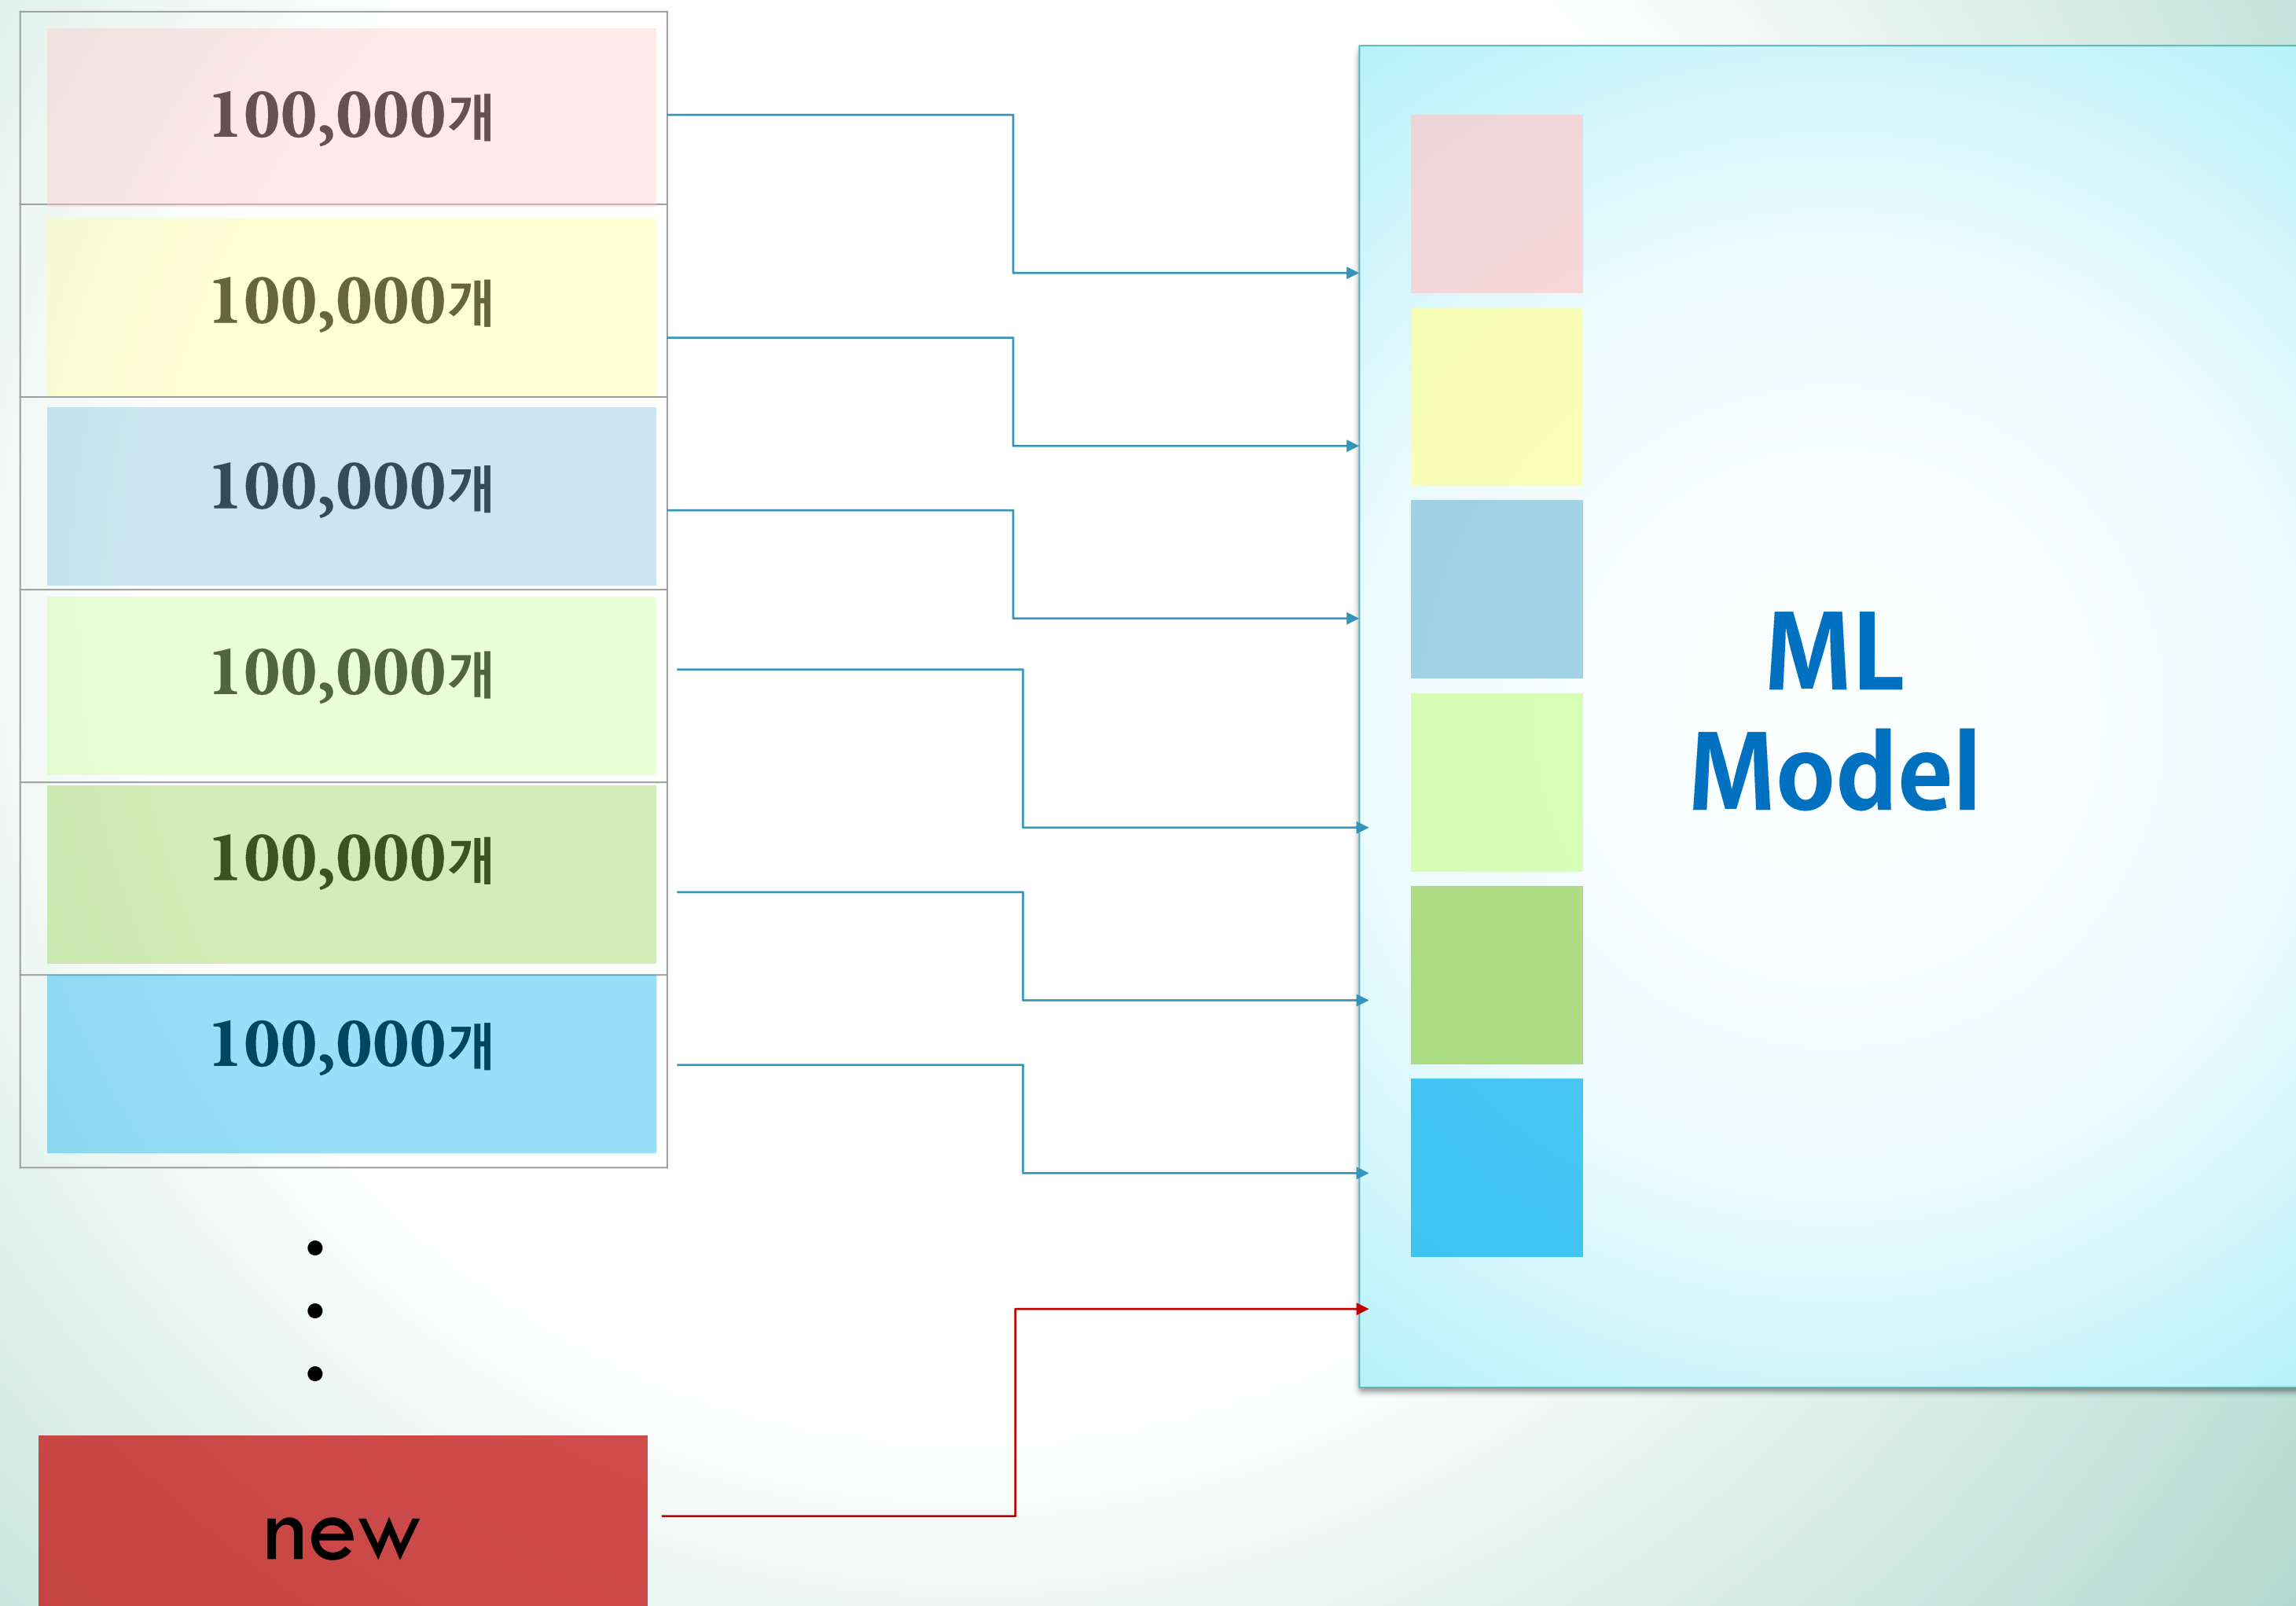
$$L = \frac{1}{N} \sum_i D(S(WX_i + b), L_i) + \lambda \sum w^2$$

Training set

The equation shows the total cost (loss) L as the sum of the average loss on the training set and a regularization term. The first term, $\frac{1}{N} \sum_i D(S(WX_i + b), L_i)$, represents the loss on the training set, where i indexes the training samples. The second term, $\lambda \sum w^2$, represents the regularization term, where λ is the regularization strength and w represents the weights. A red arrow points from the 'Training set' label to the index i in the first term.

Online learning

1,000,000개



MNIST Dataset



[train-images-idx3-ubyte.gz](#): training set images (9912422 bytes)
[train-labels-idx1-ubyte.gz](#): training set labels (28881 bytes)
[t10k-images-idx3-ubyte.gz](#): test set images (1648877 bytes)
[t10k-labels-idx1-ubyte.gz](#): test set labels (4542 bytes)

MNIST Dataset

[illegible]

https://docs.google.com/spreadsheets/d/1Phwp2oCrHz08939_T0LqMyX-kbxun96x8yq3g0Rgitl/edit#gid=0

Accuracy

- **How many of your predictions are correct?**
- **95% ~ 99%?**